Project 2: Supervised Learning

Building a Student Intervention System

1. Classification vs Regression

Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?

Answer: I choose classification model to predict 'Pass' or 'Fail' and it is easily find who needs intervention. On the other hand regression, it may predict probability of pass but I see it's hard to make threshold when we start intervension.

2. Exploring the Data

Let's go ahead and read in the student dataset first.

*To execute a code cell, click inside it and press **Shift+Enter**.*

In [139]:

```
# Import libraries
import numpy as np
import pandas as pd
```

In [140]:

```
# Read student data
student_data = pd.read_csv("student-data.csv")
print "Student data read successfully!"
# Note: The last column 'passed' is the target/label, all other are feature columns
```

```
Student data read successfully!
```

Now, can you find out the following facts about the dataset?

- Total number of students
- Number of students who passed
- Number of students who failed
- Graduation rate of the class (%)
- Number of features

*Use the code block below to compute these values. Instructions/steps are marked using **TODO**s.*

In [141]:

```
# TODO: Compute desired values - replace each '?' with an appropriate expression/function
call
n_students = len(student_data)
n_features = student_data.shape[1]-1
n_passed = len(student_data.ix[(student_data['passed']=='yes')])
n_failed = len(student_data.ix[(student_data['passed']=='no')])
grad_rate = n_passed * 100.0 / n_students
print "Total number of students: {}".format(n_students)
print "Number of students who passed: {}".format(n_passed)
print "Number of students who failed: {}".format(n_failed)
print "Number of features: {}".format(n_features)
print "Graduation rate of the class: {:.2f}%".format(grad_rate)
```

```
Total number of students: 395
Number of students who passed: 265
Number of students who failed: 130
Number of features: 30
Graduation rate of the class: 67.09%
```

3. Preparing the Data

In this section, we will prepare the data for modeling, training and testing.

Identify feature and target columns

It is often the case that the data you obtain contains non-numeric features. This can be a problem, as most machine learning algorithms expect numeric data to perform computations with.

Let's first separate our data into feature and target columns, and see if any features are non-numeric.
Note: For this dataset, the last column (`'passed'`) is the target or label we are trying to predict.

In [142]:

```
# Extract feature (X) and target (y) columns
feature_cols = list(student_data.columns[:-1])  # all columns but last are features
target_col = student_data.columns[-1]  # last column is the target/label
print "Feature column(s):-¥n{}".format(feature_cols)
print "Target column: {}".format(target_col)

X_all = student_data[feature_cols]  # feature values for all students
y_all = student_data[target_col]  # corresponding targets/labels
print "¥nFeature values:-"
print X_all.head()  # print the first 5 rows
```

```
Feature column(s):-
['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu', 'M
job', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime', 'failures',
'schoolsup', 'famsup', 'paid', 'activities', 'nursery', 'higher', 'interne
t', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc', 'Walc', 'health', 'ab
sences']
Target column: passed

Feature values:-
  school sex  age address famsize Pstatus  Medu  Fedu     Mjob      Fjob  ¥
0     GP   F   18       U     GT3       A     4     4  at_home   teacher
1     GP   F   17       U     GT3       T     1     1  at_home     other
2     GP   F   15       U     LE3       T     1     1  at_home     other
3     GP   F   15       U     GT3       T     4     2   health  services
4     GP   F   16       U     GT3       T     3     3    other     other

    ...    higher internet  romantic  famrel  freetime goout Dalc Walc healt
h  ¥
0   ...       yes       no        no       4         3     4    1    1
3
1   ...       yes      yes        no       5         3     3    1    1
3
2   ...       yes      yes        no       4         3     2    2    3
3
3   ...       yes      yes       yes       3         2     2    1    1
5
4   ...       yes       no        no       4         3     2    1    2
5

  absences
0        6
1        4
2       10
3        2
4        4

[5 rows x 30 columns]
```

Preprocess feature columns

As you can see, there are several non-numeric columns that need to be converted! Many of them are simply `yes`/`no`, e.g. `internet`. These can be reasonably converted into 1/0 (binary) values.

Other columns, like `Mjob` and `Fjob`, have more than two values, and are known as *categorical variables*. The recommended way to handle such a column is to create as many columns as possible values (e.g. `Fjob_teacher`, `Fjob_other`, `Fjob_services`, etc.), and assign a 1 to one of them and 0 to all others.

These generated columns are sometimes called *dummy variables*, and we will use the [`pandas.get_dummies()` (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.get_dummies.html?highlight=get_dummies#pandas.get_dummies)](http://pandas.pydata.org/pandas-docs/stable/generated/pandas.get_dummies.html?highlight=get_dummies#pandas.get_dummies) function to perform this transformation.

In [143]:

```
# Preprocess feature columns
def preprocess_features(X):
    outX = pd.DataFrame(index=X.index)  # output dataframe, initially empty

    # Check each column
    for col, col_data in X.iteritems():
        # If data type is non-numeric, try to replace all yes/no values with 1/0
        if col_data.dtype == object:
            col_data = col_data.replace(['yes', 'no'], [1, 0])
        # Note: This should change the data type for yes/no columns to int

        # If still non-numeric, convert to one or more dummy variables
        if col_data.dtype == object:
            col_data = pd.get_dummies(col_data, prefix=col)  # e.g. 'school' => 'school_G
P', 'school_MS'

        outX = outX.join(col_data)  # collect column(s) in output dataframe

    return outX

X_all = preprocess_features(X_all)
print "Processed feature columns ({}):-¥n{}".format(len(X_all.columns), list(X_all.column
s))
```

```
Processed feature columns (48):-
['school_GP', 'school_MS', 'sex_F', 'sex_M', 'age', 'address_R', 'address_
U', 'famsize_GT3', 'famsize_LE3', 'Pstatus_A', 'Pstatus_T', 'Medu', 'Fedu',
'Mjob_at_home', 'Mjob_health', 'Mjob_other', 'Mjob_services', 'Mjob_teache
r', 'Fjob_at_home', 'Fjob_health', 'Fjob_other', 'Fjob_services', 'Fjob_teac
her', 'reason_course', 'reason_home', 'reason_other', 'reason_reputation',
'guardian_father', 'guardian_mother', 'guardian_other', 'traveltime', 'study
time', 'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc', 'Wa
lc', 'health', 'absences']
```

Split data into training and test sets

So far, we have converted all *categorical* features into numeric values. In this next step, we split the data (both features and corresponding labels) into training and test sets.

In [144]:

```
# First, decide how many training vs test samples you want
num_all = student_data.shape[0]  # same as len(student_data)
num_train = 300  # about 75% of the data
num_test = num_all - num_train

# TODO: Then, select features (X) and corresponding labels (y) for the training and test s
ets
# Note: Shuffle the data or randomly select samples to avoid any bias due to ordering in t
he dataset
from sklearn.cross_validation import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X_all,y_all,test_size=num_test,random_state
=0)
print "Training set: {} samples".format(X_train.shape[0])
print "Test set: {} samples".format(X_test.shape[0])
# Note: If you need a validation set, extract it from within training data
```

```
Training set: 300 samples
Test set: 95 samples
```

4. Training and Evaluating Models

Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem. For each model:

- What are the general applications of this model? What are its strengths and weaknesses?
- Given what you know about the data so far, why did you choose this model to apply?
- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F_1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.

Produce a table showing training time, prediction time, F_1 score on training set and F_1 score on test set, for each training set size.

Note: You need to produce 3 such tables - one for each model.

In [145]:

```
# Train a model
import time

def train_classifier(clf, X_train, y_train):
    #print "Training {}...".format(clf.__class__.__name__)
    start = time.time()
    clf.fit(X_train, y_train)
    end = time.time()
    #print "|TrainTime:{:.3f}|".format(end - start)
    return end-start
    #print "Done!¥nTraining time (secs): {:.3f}".format(end - start)

# TODO: Choose a model, import it and instantiate an object
from sklearn.neighbors import NearestNeighbors
from sklearn.svm import SVC
clf = SVC()

# Fit model to training data
train_classifier(clf, X_train, y_train)  # note: using entire training set here
print clf  # you can inspect the learned model by printing it
```

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

In [146]:

```
# Predict on training set and compute F1 score
from sklearn.metrics import f1_score

def predict_labels(clf, features, target):
    #print "Predicting labels using {}...".format(clf.__class__.__name__)
    start = time.time()
    y_pred = clf.predict(features)
    end = time.time()
    #print "Done!¥nPrediction time (secs): {:.3f}".format(end - start)
    #print "|PredTime{:.3f}|".format(end - start)
    return [(end-start),f1_score(target.values, y_pred, pos_label='yes')]

train_f1_score = predict_labels(clf, X_train, y_train)
print "F1 score for training set: {}".format(train_f1_score)
```

```
F1 score for training set: [0.006000041961669922, 0.86919831223628696]
```

In [147]:

```
# Predict on test data
print "F1 score for test set: {}".format(predict_labels(clf, X_test, y_test))
```

```
F1 score for test set: [0.003000020980834961, 0.75862068965517238]
```

In [148]:

```
# Train and predict using different training set sizes
def train_predict(clf, X_train, y_train, X_test, y_test):
    #print "------------------------------------------"
    #print "Training set size: {}".format(len(X_train))
    train_time=train_classifier(clf, X_train, y_train)
    #print "F1 score for training set: {}".format(predict_labels(clf, X_train, y_train))
    #print "F1 score for test set: {}".format(predict_labels(clf, X_test, y_test))
    train_result=predict_labels(clf, X_train, y_train)
    test_result=predict_labels(clf, X_test, y_test)
    print "| {} | {:.3f} | {:.3f} | {:.3f}| {:.5f} | {:.5f} |".format(len(X_train) , train
_time, train_result[0],test_result[0],train_result[1],test_result[1] )



# TODO: Run the helper function above for desired subsets of training data
clf =SVC()
print " | {} | Training Time | Train Pred Time | Test Pred Time | F1_Train | F1_Test |".fo
rmat(clf.__class__.__name__)
print "  | --- | -- | -- | --| -- | -- |"
for x, i in enumerate([100,200,300]):
    train_predict(clf,X_train[:i],y_train[:i],X_test,y_test)
# Note: Keep the test set constant
```

```
 | SVC | Training Time | Train Pred Time | Test Pred Time | F1_Train | F1_Te
st |
  | --- | -- | -- | --| -- | -- |
| 100 | 0.001 | 0.001 | 0.001| 0.85906 | 0.78378 |
| 200 | 0.007 | 0.003 | 0.002| 0.86928 | 0.77551 |
| 300 | 0.009 | 0.008 | 0.002| 0.86920 | 0.75862 |
```

In [ ]:

In [165]:

```
# TODO: Train and predict using two other models
#SGD
from sklearn.linear_model import SGDClassifier
clf2 = SGDClassifier(loss="hinge", penalty="l2")
print " | {} | Training Time | Train Pred Time | Test Pred Time | F1_Train | F1_Test |".fo
rmat(clf2.__class__.__name__)
print "  | --- | -- | -- | --| -- | -- |"
for x, i in enumerate([100,200,300]):
    train_predict(clf2,X_train[:i],y_train[:i],X_test,y_test)

#Neighbors
from sklearn import neighbors
clf3 = neighbors.KNeighborsClassifier()
print " | {} | Training Time | Train Pred Time | Test Pred Time | F1_Train | F1_Test |".fo
rmat(clf3.__class__.__name__)
print "  | --- | -- | -- | --| -- | -- |"
for x, i in enumerate([100,200,300]):
    train_predict(clf3,X_train[:i],y_train[:i],X_test,y_test)
```

```
 | SGDClassifier | Training Time | Train Pred Time | Test Pred Time | F1_Tra
in | F1_Test |
  | --- | -- | -- | --| -- | -- |
| 100 | 0.002 | 0.004 | 0.015| 0.67257 | 0.68376 |
| 200 | 0.001 | 0.016 | 0.000| 0.79880 | 0.77124 |
| 300 | 0.002 | 0.000 | 0.000| 0.81600 | 0.78146 |
 | KNeighborsClassifier | Training Time | Train Pred Time | Test Pred Time |
F1_Train | F1_Test |
  | --- | -- | -- | --| -- | -- |
| 100 | 0.001 | 0.001 | 0.002| 0.79720 | 0.70677 |
| 200 | 0.004 | 0.006 | 0.010| 0.85714 | 0.71212 |
| 300 | 0.002 | 0.012 | 0.005| 0.87225 | 0.74820 |
```

Result:

(All time are in seconds)

SVC-Training Set	Training Time	Train Pred Time	Test Pred Time	F1_Train	F1_Test
100	0.001	0.001	0.001	0.85906	0.78378
200	0.003	0.003	0.002	0.86928	0.77551
300	0.008	0.005	0.002	0.86920	0.75862

SGDClassifier -Training Set	Training Time	Train Pred Time	Test Pred Time	F1_Train	F1_Test
100	0.002	0.004	0.015	0.67257	0.68376
200	0.001	0.016	0.000	0.79880	0.77124
300	0.002	0.000	0.000	0.81600	0.78146

KNeighborsClassifier - Training Set	Training Time	Train Pred Time	Test Pred Time	F1_Train	F1_Test
100	0.001	0.002	0.004	0.79720	0.70677
200	0.000	0.005	0.002	0.85714	0.71212
300	0.001	0.019	0.007	0.87225	0.74820

5. Choosing the Best Model

- Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?
- In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a Decision Tree or Support Vector Machine, how does it make a prediction).
- Fine-tune the model. Use Gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.
- What is the model's final F_1 score?

Answer

- Comparing 3 models of SVC,SGD Classifier and KNeighborsClassifiler, I recommend SGD Classifier because of 2 major reasons. 1) Training Time doesn't increase as training sets increas. 2)F1 scores for Testing is relatively high rather than others.'
- SGD (stochastic gradient descent) works as 1) randomly picks up few dataset from training data, 2) predict and get the difference from right answer, 3) reflect the difference to optimize the model , 4) repeat 1)-3) process till the model gets minimum difference.

In [151]:

```
# TODO: Fine-tune your model and report the best F1 score
```

In [180]:

```
from sklearn.grid_search import GridSearchCV
print clf2
clf4 = SGDClassifier(loss="hinge")
parameters = {'alpha':(0.1,0.001,0.0001,0.00001),'loss':('hinge','log','modified_huber','s
quared_hinge','perceptron'),'penalty':('none','l1','elasticnet','l2')}
#scoring_function = make_scorer(mean_absolute_error,greater_is_better=False)
reg = GridSearchCV(clf4,parameters)

print " | {} | Training Time | Train Pred Time | Test Pred Time | F1_Train | F1_Test |".fo
rmat(reg.__class__.__name__)
print "  | --- | -- | -- | --| -- | -- |"
train_predict(reg,X_train,y_train,X_test,y_test)

print "Final model optimal parameters:", reg.best_params_
```

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None, epsilon=0.1,
       eta0=0.0, fit_intercept=True, l1_ratio=0.15,
       learning_rate='optimal', loss='hinge', n_iter=5, n_jobs=1,
       penalty='l2', power_t=0.5, random_state=None, shuffle=True,
       verbose=0, warm_start=False)
 | GridSearchCV | Training Time | Train Pred Time | Test Pred Time | F1_Trai
n | F1_Test |
  | --- | -- | -- | --| -- | -- |
| 300 | 1.370 | 0.000 | 0.001| 0.82186 | 0.78378 |
Final model optimal parameters: {'penalty': 'elasticnet', 'alpha': 0.1, 'los
s': 'hinge'}
```

Finetune result

GridSearchCV found the best parameters as 'penalty'='elasticnet', 'alpha'= 0.1, 'loss'= 'hinge' for SGDClassifier.

GridSearchCV - Training Set	Training Time	Train Pred Time	Test Pred Time	F1_Train	F1_Test
300	1.370	0.000	0.001	0.82186	0.78378

this F1_test score is slightly better than 0.78146 with default setup.